
Transformixer: Attention-Based Variate Mixing without Positional Bias for Multivariate Forecasting

Morteza Ziaabakhsh
University of Tehran
m.ziaabakhsh@ut.ac.ir

Abstract

Multivariate long-term forecasting requires models that capture both temporal dynamics and cross-variate dependencies. Recent mixer architectures such as xLSTM-Mixer combine an efficient shared linear temporal forecast with an xLSTM-based joint variate mixer. However, recurrent mixing along the variate axis is inherently order-sensitive: it depends on an arbitrary channel ordering, propagates information sequentially across channels, and relies on a forward/reverse multi-view mechanism to mitigate that dependence. We propose **Transformixer**, which retains reversible instance normalization and shared NLinear temporal mixing, but replaces the ordered sLSTM variate mixer with a Transformer encoder over variate tokens *without* positional encodings. The resulting mixer is permutation-equivariant, realizes direct pairwise interactions among all variates in a single layer, and removes the need for memory tokens and bidirectional multi-view fusion. On Electricity and Traffic with lookback and horizon 96, Transformixer is competitive with strong baselines: it improves upon xLSTM-Mixer on Traffic (including the best MAE among compared models) while remaining slightly behind on Electricity. Our study is intentionally limited in breadth by available compute and time (two datasets, one horizon); nevertheless, the present results suggest that Transformixer is a promising direction, and future work can explore this idea more thoroughly across additional datasets, horizons, and training budgets.

1 Introduction

Multivariate time series forecasting is central to energy, traffic, and industrial monitoring, where many correlated channels are observed jointly. Deep models must therefore learn temporal structure *and* cross-channel interactions. Transformers [1] have been widely adapted to this setting, from efficient temporal attention [2] and patch-based channel-independent encoders [3] to inverted architectures that attend over variate tokens [4]. In parallel, strong linear and mixer baselines [5–7] show that careful temporal mapping and lightweight mixing remain highly competitive.

xLSTM-Mixer [7] is representative of the latter line: a shared NLinear temporal forecast is refined by an xLSTM [8] (sLSTM) block that mixes information along the *variate* axis, optionally with memory tokens, and reconciles forward and reversed channel views. This design is effective, but it encodes a questionable inductive bias whenever channel indices lack a natural total order (e.g., sensors or customer meters). Sequential mixing makes the computation depend on an arbitrary ordering, forces long-range channel interactions to travel through a recurrent state, and motivates an auxiliary reverse view to reduce order artifacts.

We ask whether a simpler, set-structured mixer is a better fit. **Transformixer** keeps the successful temporal front-end of xLSTM-Mixer, reversible instance normalization [9] and shared NLinear temporal mixing [5], but replaces the sLSTM variate mixer with a Transformer encoder applied to variate tokens *without* positional encodings and without a causal mask. Beyond permutation

equivariance [10], attention provides direct pairwise interactions among all channels, multi-head specialization, and a single forward pass that no longer needs memory tokens or bidirectional view fusion. This view is complementary to iTransformer [4]: we reuse variate-token attention as a *mixer block* after an explicit temporal linear forecast, rather than as a full inverted backbone.

Contributions.

- We introduce Transformixer, a drop-in permutation-equivariant Transformer variate mixer for the xLSTM-Mixer pipeline.
- We analyze inductive-bias differences versus ordered sLSTM mixing (equivariance, global interactions, architectural simplification, and parallelism), not only forecasting accuracy.
- We compare against five baselines on Electricity and Traffic ($L=H=96$) and discuss when the redesign helps, under clearly stated limitations.

2 Related Work

Transformer architectures for long-term forecasting. The Transformer [1] has become a dominant backbone for long-term time series forecasting (LTSF), motivating a line of work that adapts attention to long horizons and non-stationary signals. Informer [2] reduces the cost of canonical self-attention with ProbSparse attention and a generative decoder, enabling longer lookbacks than early encoder-decoder designs. Autoformer [11] embeds seasonal-trend decomposition into the network and replaces standard attention with an auto-correlation mechanism that emphasizes period-based dependencies. FEDformer [12] further couples decomposition with frequency-domain attention, arguing that sparse spectral representations capture global structure more effectively than purely temporal attention. Subsequent models broaden the representational toolkit: TimesNet [13] reshapes 1D series into 2D tensors to model intra- and inter-period variation, while PatchTST [3] segments each univariate channel into patches and applies a channel-independent Transformer, often matching or exceeding heavier multivariate attention models on standard LTSF benchmarks. These works primarily innovate on *temporal* representation and efficiency; cross-variate interaction, when present, is typically entangled with time-axis attention rather than isolated as an explicit mixer.

Channel independence versus cross-variate modeling. A recurring design tension in multivariate forecasting is whether channels should be modeled independently or jointly. Channel-independent approaches such as PatchTST [3] and strong linear baselines [5, 6] share temporal parameters across variates but avoid explicit cross-channel mixing, which can improve robustness when inter-variate relations are weak or unstable. In contrast, several Transformers target cross-dimension dependency directly. Crossformer [14] uses dimension-segment-wise embeddings and two-stage attention to capture cross-time and cross-dimension interactions hierarchically. iTransformer [4] inverts the conventional layout: each variate’s full series becomes a token, so attention mixes channels while feed-forward layers process series representations. This inverted view is closely related to Transformixer’s variate-token mixer; the difference is architectural placement. iTransformer is a full inverted backbone, whereas Transformixer applies PE-free attention as a *drop-in variate mixer* after an explicit shared NLinear temporal forecast inside a mixer-style pipeline.

Linear models and MLP mixers. Challenging the necessity of deep temporal attention, LTSF-Linear / DLinear / NLinear [5] show that simple linear maps, especially with last-value normalization, can rival Transformers on many benchmarks. RLinear [6] revisits linear mapping with RevIN-style normalization and reinforces that carefully normalized linear temporal projections remain highly competitive. Building on MLP-Mixer ideas from vision, TSMixer [15] alternates time-mixing and feature-mixing MLP blocks to capture temporal and cross-variate structure with favorable parameter scaling. TimeMixer [16] extends mixing to multiple temporal resolutions via past decomposable mixing and future multipredictor mixing. Collectively, these results motivate hybrid designs that keep an efficient temporal front-end and allocate capacity where cross-variate reasoning is needed, the regime in which xLSTM-Mixer and Transformixer operate.

Recurrent mixers and xLSTM. Despite the Transformer wave, recurrent models have regained attention through xLSTM [8], which revisits LSTM-style memory with exponential gating and enriched memory structures (including scalar sLSTM blocks). xLSTM-Mixer [7] instantiates this




Figure 1: Transformixer overview. Temporal structure is handled by shared NLinear; cross-variate structure is handled by a PE-free Transformer over variate tokens.

idea for multivariate forecasting: a shared linear (NLinear) temporal forecast is refined by sLSTM blocks that mix along the variate axis, optionally with memory tokens, and reconcile forward and reversed channel views. The design demonstrates that recurrent joint mixing can be state-of-the-art, but it also inherits order sensitivity along channels and relies on multi-view fusion to mitigate arbitrary variate orderings. Transformixer keeps the RevIN + NLinear temporal pathway of this lineage while replacing the ordered sLSTM mixer with set-structured self-attention.

Normalization and set inductive biases. Non-stationarity is a central obstacle in LTSF. RevIN [9] applies reversible instance normalization to reduce distribution shift between train and test windows and is now a standard component in many competitive models, including both xLSTM-Mixer and Transformixer. Separately, set learning theory [10] formalizes permutation-invariant and equivariant maps over unordered collections. When multivariate channels lack a natural total order, as with sensors, meters, or traffic detectors, treating variates as a set is a natural inductive bias. Transformixer’s omission of positional encodings along the variate axis is a concrete realization of that principle: the mixer is permutation-equivariant by construction, in contrast to recurrent ordered mixing and to Transformer mixers that inject positional signals on the channel axis.

Positioning. In summary, prior work offers strong temporal Transformers, channel-independent patch models, linear and MLP mixers, inverted variate-token attention, and recurrent joint mixers. Transformixer sits at their intersection: it inherits an efficient linear temporal forecast from the mixer/linear literature, adopts variate-token attention in the spirit of iTransformer, and specifically targets the order-sensitive bottleneck of xLSTM-Mixer’s sLSTM variate stage with a PE-free, single-view Transformer mixer.

3 Method

3.1 Problem setup

Given a lookback window $X \in \mathbb{R}^{L \times C}$ with C variates, the goal is to predict the future horizon $Y \in \mathbb{R}^{H \times C}$. We consider multivariate-to-multivariate forecasting. Transformixer implements a map $f_\theta : \mathbb{R}^{L \times C} \rightarrow \mathbb{R}^{H \times C}$ composed of normalization, temporal mixing, equivariant variate mixing, and denormalization (Figure 1).

3.2 Reversible instance normalization

Following Kim et al. [9], we apply RevIN with non-affine statistics along the temporal axis of each variate, normalize the lookback window, and denormalize the forecast at the end of the network. This reduces sensitivity to non-stationarity without altering the mixer design.

3.3 Shared NLinear temporal mixing

Let $x \in \mathbb{R}^{B \times L \times C}$ denote the RevIN-normalized batch. We use the NLinear residual map of Zeng et al. [5]: subtract the last observed value, apply a linear layer shared across channels that maps length L to horizon H , then add the last value back,

$$\hat{Y}_0 = \text{Linear}_{L \rightarrow H}(x - x_{:, -1, :}) + x_{:, -1, :}, \quad (1)$$

yielding a preliminary forecast $\hat{Y}_0 \in \mathbb{R}^{B \times H \times C}$. This stage models temporal dynamics efficiently and independently of channel order.

Table 1: Transformixer versus xLSTM-Mixer (FULL setting used in our runs).

Component	xLSTM-Mixer	Transformixer
RevIN	yes	yes
Shared NLinear temporal mix	yes	yes
Variate mixer	sLSTM (+ mem. tokens)	Transformer encoder
Positional encoding on variates	n/a (order via recurrence)	none
Forward/reverse multi-view	yes	no
Memory tokens	yes (3 in our runs)	no

3.4 Variate tokenization and Transformer mixing

We rearrange $\hat{Y}_0$ so that each variate becomes a token of dimension H , embed tokens with a linear map $H \rightarrow d$, and process the sequence of C tokens with a stack of standard Transformer encoder layers [1]: multi-head self-attention and a position-wise feed-forward network, with residual connections and layer normalization. Crucially, we *omit positional encodings* and use full (non-causal) attention. A final linear map $d \rightarrow H$ produces refined horizon vectors per variate, which are rearranged and RevIN-denormalized to obtain $\hat{Y}$.

Permutation equivariance. Because the embedding, attention, feed-forward maps, and output projection are applied identically to every token and attention is a symmetric function of the token set, for any channel permutation matrix P ,

$$f_{\theta}(XP) = f_{\theta}(X)P. \quad (2)$$

Thus the model treats variates as a set rather than an ordered sequence. We verify this property with a dedicated unit test in our implementation.

3.5 Design rationale versus xLSTM-Mixer

Table 1 summarizes what Transformixer keeps and removes. Beyond equivariance, replacing sLSTM with attention is motivated by several inductive biases:

1. **Direct global interactions.** Self-attention couples every pair of variates in one layer ($O(1)$ path length over channels). An sLSTM mixes channels sequentially, so information from channel i to j must travel through the recurrent state (and/or a small set of memory tokens), with path length depending on an arbitrary ordering. This distinction is especially relevant for high- C series such as Traffic ($C=862$).
2. **Multi-head specialization.** Different heads can learn distinct dependency patterns in parallel, whereas a single recurrent state stream is less naturally factored into multiple interaction types.
3. **Architectural simplification.** Order-agnostic mixing removes the need for xLSTM-Mixer’s forward/reverse multi-view fusion and for auxiliary memory tokens. Transformixer uses a single forward pass over variate tokens.
4. **Hardware-parallel mixing.** Attention is token-parallel over C , whereas recurrence is sequential in C . We do *not* claim wall-clock superiority: attention costs $O(C^2d)$ per layer and may become expensive for very large C .

4 Experimental Setup

Datasets. We evaluate on two standard multivariate benchmarks commonly used in long-term forecasting [2, 11]: **Electricity** (ECL; $C=321$ hourly client loads) and **Traffic** ($C=862$ hourly road occupancy rates). Following common practice, we use a chronological 70%/10%/20% train/validation/test split and fit a StandardScaler on the training portion only.

Table 2: Forecasting results for lookback/horizon 96. Best is **bold**; second-best is underlined.

Model	ECL		Traffic	
	MSE	MAE	MSE	MAE
RLinear	0.1973	0.2740	0.6463	0.3844
Informer	0.3263	0.4068	1.0140	0.5624
PatchTST	0.1754	0.2603	0.4699	0.3005
iTransformer	<u>0.1481</u>	0.2398	0.3931	0.2687
xLSTM-Mixer	0.1480	0.2352	0.4157	<u>0.2607</u>
Transformixer	0.1517	<u>0.2364</u>	<u>0.4068</u>	0.2558

Protocol. All models use lookback $L=96$ and prediction horizon $H=96$ in the multivariate setting. We report mean squared error (MSE) and mean absolute error (MAE) on the denormalized forecasts. We compare Transformixer to five baselines that we re-ran in the same project setting: RLinear [6], Informer [2], PatchTST [3], iTransformer [4], and xLSTM-Mixer [7].

Training details. Transformixer and xLSTM-Mixer are trained with Adam, learning rate 5×10^{-4} , batch size 32, L1 loss, and up to 10 epochs with a warmup–cosine schedule (seed 42). Both use embedding width 1024, 16 heads, and 2 mixer blocks/layers with dropout 0.1. For each baseline, we follow the hyperparameters recommended by the authors of the corresponding paper (or the accompanying official codebase) for the given dataset setting, including model width, depth, learning rate, and related training choices where applicable. In contrast, Transformixer does *not* benefit from a dedicated hyperparameter search: its architectural and optimization settings are largely inherited from the xLSTM-Mixer experimental scripts rather than tuned specifically for the Transformer mixer. The baselines therefore enjoy published, author-recommended configurations, while Transformixer is evaluated under a comparatively untuned setup. We treat the study as an architectural comparison under practical training budgets rather than a fully unified or exhaustively tuned hyperparameter sweep.

Scope limitations. Due to limited computational resources and project time, we report a single run per model–dataset pair, only horizon 96, and only two datasets. We do not claim comprehensive state-of-the-art coverage; broader multi-horizon and multi-seed evaluation is left to future work (Section 6).

5 Results

Table 2 and Figure 2 summarize the main comparison. Transformixer is competitive overall and does not dominate every cell.

On **Traffic**, Transformixer improves upon its direct predecessor xLSTM-Mixer on both MSE (0.4068 vs. 0.4157) and MAE (0.2558 vs. 0.2607), achieving the best MAE among all compared models and the second-best MSE after iTransformer. This is consistent with the hypothesis that direct, order-agnostic attention is particularly useful when many channels interact without a meaningful 1-D order.

On **ECL**, xLSTM-Mixer and iTransformer remain slightly ahead (MSE 0.1480 / 0.1481 versus 0.1517 for Transformixer). The gap is modest ($\approx 2.5\%$ relative MSE versus xLSTM-Mixer), suggesting that ordered sLSTM mixing with memory tokens remains a strong alternative at moderate C , or that Transformixer hyperparameters inherited from the xLSTM-Mixer setup are not fully optimal for ECL.

Weaker baselines (Informer, RLinear) lag substantially on both datasets, indicating that the evaluation setting is non-trivial and that the stronger inverted/mixer models form a meaningful comparison group. Relative to xLSTM-Mixer specifically, Transformixer trades a small regression on ECL for a clear gain on the higher-dimensional Traffic benchmark (Figure 3).




Figure 2: MSE (top) and MAE (bottom) on ECL and Traffic ($L=H=96$). Transformer bars are outlined.

5.1 Permutation behavior

By construction, Transformer is permutation-equivariant along channels: for a channel permutation matrix P , $f(XP) = f(X)P$. Unit tests in our implementation confirm this identity up to numerical tolerance. In contrast, xLSTM-Mixer’s sLSTM mixer consumes an ordered sequence of variate tokens (with an optional reverse view), so the computation itself depends on channel order.

To probe the practical effect of this distinction, we load the Traffic-trained Transformer and xLSTM-Mixer checkpoints and evaluate MSE/MAE under the identity channel order and $K=8$ independent random permutations of the variate axis (Table 3). Transformer’s errors are essentially unchanged across permutations (MSE std 1.4×10^{-3}), as expected from equivariance. xLSTM-Mixer shows similarly small metric drift (MSE std 1.5×10^{-3}), consistent with forward/reverse multi-view partially mitigating order artifacts, but remains strictly worse on every run: mean MSE 1.349 vs. 1.286 ($\approx 4.6\%$ relative) and mean MAE 0.923 vs. 0.903 ($\approx 2.2\%$ relative). The table makes the comparison explicit: both models are fairly stable under channel shuffles in this probe, yet Transformer occupies a clearly better error regime without relying on ordered mixing.

Figure 4 visualizes a Transformer attention map (first head, 64-variate subset) from the same Traffic checkpoint. The weights are sparse and non-uniform: a small number of key variates attract attention from many queries, indicating structured cross-variate interactions rather than uniform or purely diagonal self-attention. This qualitative pattern aligns with the claim that the mixer models direct pairwise dependencies among channels.




Figure 3: Relative MSE change of Transformer versus xLSTM-Mixer (negative is better for Transformer).

Table 3: Permutation probe with Traffic-trained checkpoints: identity order versus summary over the identity and $K=8$ random channel permutations. Lower is better. Transformer is nearly invariant and consistently more accurate.

Model	Identity		Over 9 orders (mean $\pm$ std)	
	MSE	MAE	MSE	MAE
xLSTM-Mixer	1.3463	0.9224	1.3486 $\pm$ 0.0015	0.9233 $\pm$ 0.0005
Transformer	1.2841	0.9019	1.2865 $\pm$ 0.0014	0.9027 $\pm$ 0.0005

6 Discussion

When does the Transformer mixer help? The Traffic improvement aligns with the motivation in Section 3.5: high- C forecasting benefits from direct pairwise mixing and set equivariance. The permutation probe in Section 5.1 further shows that Transformer remains stable under channel shuffles and consistently outperforms xLSTM-Mixer on that controlled evaluation, while the attention map illustrates structured cross-variate interactions. On ECL, the slightly better xLSTM-Mixer result cautions against a universal claim that attention is a strictly superior variate mixer. Recurrent mixing with memory tokens may still be competitive when C is moderate, when training budgets are short, or when hyperparameters are tuned for the sLSTM design.

Complexity tradeoff. Attention is more hardware-parallel over channels but asymptotically quadratic in C . We did not profile latency or FLOPs in this study; efficiency claims would require dedicated measurement. For extremely large C , sparse or linearized attention could be a natural extension.

Threats to validity. Results are based on single seeds, one horizon, and two datasets. Training recipes are not fully homogenized across baseline families (loss, early stopping, seeds). More importantly for fairness, the baselines are run with author-recommended hyperparameters, whereas Transformer was not independently tuned and largely reuses xLSTM-Mixer settings. This asymmetry may understate Transformer’s potential relative to well-tuned baselines. These factors limit statistical strength and strict comparability; we report them explicitly because the project was constrained by compute and time.

Future work. A fuller evaluation would include horizons $\{96, 192, 336, 720\}$, additional datasets (e.g., ETT, Weather), multi-seed mean $\pm$ std, controlled ablations (with vs. without positional encodings; number of layers/heads), and efficiency profiling. Such experiments are feasible with additional GPU resources but were out of scope here.




Figure 4: Transformer attention weights (head 0) on a 64-variate subset from a Traffic-trained model. Bright columns indicate key variates attended to by many queries.

7 Conclusion

Transformixer replaces the order-sensitive sLSTM variate mixer in xLSTM-Mixer with a PE-free Transformer over variate tokens, preserving RevIN and shared NLinear temporal mixing. The redesign targets a more natural inductive bias for unordered channels: permutation equivariance, direct multi-head interactions, and a simpler single-view architecture without memory tokens. Empirically, Transformixer is competitive with strong baselines under a limited protocol, improving on xLSTM-Mixer for Traffic while remaining close on Electricity. We view the contribution as an architectural study with transparent resource limits, and as evidence that set-structured attention is a promising drop-in alternative for joint variate mixing in mixer-style forecasters.

Acknowledgments and Disclosure of Funding

This work was conducted as a university research project under constrained computational resources and time. Experiments were run using cloud GPU notebooks (Google Colab). We thank the authors of the baseline codebases and benchmark datasets used in this study.

References

- [1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- [2] Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 11106–11115, 2021.
- [3] Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *International Conference on Learning Representations*, 2023.
- [4] Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. iTransformer: Inverted transformers are effective for time series forecasting. In *International Conference on Learning Representations*, 2024.
- [5] Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 11121–11128, 2023.

- [6] Zhe Li, Shiyi Qi, Yiduo Li, and Zenglin Xu. Revisiting long-term time series forecasting: An investigation on linear mapping. *arXiv preprint arXiv:2305.10721*, 2023.
- [7] Maurice Kraus, Felix Divo, Devendra Singh Dhami, and Kristian Kersting. xLSTM-Mixer: Multivariate time series forecasting by mixing via scalar memories. *arXiv preprint arXiv:2410.16928*, 2024.
- [8] Maximilian Beck, Korbinian Pöppel, Markus Spanring, Andreas Auer, Oleksandra Prudnikova, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. xLSTM: Extended long short-term memory. In *Advances in Neural Information Processing Systems*, 2024.
- [9] Taesung Kim, Jinhee Kim, Yunwon Tae, Cheonbok Park, Jang-Ho Choi, and Jaegul Choo. Reversible instance normalization for accurate time-series forecasting against distribution shift. In *International Conference on Learning Representations*, 2022.
- [10] Manzil Zaheer, Satwik Kottur, Siamak Ravanbakhsh, Barnabas Poczos, Russ R Salakhutdinov, and Alexander J Smola. Deep sets. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- [11] Haixu Wu, Jiehui Xu, Jianmin Wang, and Mingsheng Long. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. In *Advances in Neural Information Processing Systems*, volume 34, pages 22419–22430, 2021.
- [12] Tian Zhou, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin. FEDformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *Proceedings of the 39th International Conference on Machine Learning*, pages 27268–27286, 2022.
- [13] Haixu Wu, Tengge Hu, Yong Liu, Hang Zhou, Jianmin Wang, and Mingsheng Long. TimesNet: Temporal 2d-variation modeling for general time series analysis. In *International Conference on Learning Representations*, 2023.
- [14] Yunhao Zhang and Junchi Yan. Crossformer: Transformer utilizing cross-dimension dependency for multivariate time series forecasting. In *International Conference on Learning Representations*, 2023.
- [15] Si-An Chen, Chun-Liang Li, Nate Yoder, Sercan O. Arik, and Tomas Pfister. TSMixer: An all-MLP architecture for time series forecasting. *Transactions on Machine Learning Research*, 2023.
- [16] Shiyu Wang, Haixu Wu, Xiaoming Shi, Tengge Hu, Huakun Luo, Lintao Ma, James Y. Zhang, and Jun Zhou. TimeMixer: Decomposable multiscale mixing for time series forecasting. In *International Conference on Learning Representations*, 2024.